

Today

1. Go over automation homework
2. Exploratory data analysis lecture
3. Visualization tutorial
4. Break
5. Group exercise – wrangling & EDA
working in small groups
6. Reconvene to go over exercise

Workflow assignment reminders

- Everyone: Self-critique paper turned in before class on Canvas
- Next week's presenters
 - Tina, Michael, Ziyi, Antonio, Hea Jung, Michelle, Jun
 - Email PDF of slides to me before class
 - For presentation, you will share screen and have 8 minutes max to present
 - I will warn you at 7 minutes to wrap it up

Syllabus updates

Week 06, 02/10: Custom Functions Part 1

- **WORKFLOW PRESENTATIONS GROUP 1**
- Finding new packages/APIs vs. writing your own functions
- Defining custom functions within a script
- *SKILLS*: Writing basic functions

Week 07, 02/17: Custom Functions Part 2

- **WORKFLOW PRESENTATIONS GROUP 2**
- “Technical debt”, “design smells”, and code refactoring
- Sourcing functions
- Working with function arguments
- *SKILLS*: Writing advanced functions

OTHER READING: Technical debt (Suryanarayana, Samarthiyam, & Sharma, 2014)

Week 08, 02/24: Data Sharing and Reproducibility

- Reuse-minded project management
- Reproducible reports
- Preserving programming environment and analyses
- *SKILLS*: R Markdown, package control

R4DS: [R Markdown](#)

OTHER READINGS: Transparency in psychological science (Klein et al., 2018) and Care and feeding of data (Goodman et al., 2014)

Week 09, 03/03: Communication

- Communicating through graphical styles
- Interactive plots for data exploration
- Manuscript preparation in R Markdown
- *SKILLS*: ggplot and extensions, papaja

R4DS: [Graphics for communication](#)

OTHER READING: Designing graphs for decision-makers (Zacks & Franconeri, 2020), [Chartjunk](#) from (Tufte, 1990, 2001, 2006), and (optional) Graph construction (Witt, 2019).

GitHub updates

- Adding any code Jake writes when going over homework
- Tagging repos so that they're a bit easier to find

259-wrangling-homework

homework week3 wrangling

● R 2 0 0 0 Updated 20 hours ago

259-syllabus-readings-slides

Syllabus and readings

slides readings syllabus week1 assignments

● TeX 2 0 0 0 Updated 3 days ago

259-data-wrangling

tutorial week3 wrangling

● R 0 0 0 0 Updated 7 days ago

259-automation

automation tutorials

automation tutorial week4

● R 1 0 0 0 Updated 7 days ago

259-tidying-automation-homework

homework week4

● R 11 0 0 0 Updated 7 days ago

PSYC 259:

Principles of Data Science

Week 5: Exploratory Data Analysis
Error Checking & Visualization

Exploratory Data Analysis

CDA vs. EDA





- Goals of Confirmatory Data Analysis
 - Hypothesis testing, probabilistic modeling, inference
- Goals of Exploratory Data Analysis (Tukey)
 - Understanding the patterns in the data
 - Generating hypotheses
 - Checking your assumptions about data quality
 - "To find the unexpected, to avoid being fooled, and to develop rich descriptions" (Behrens & Yu, 2003)

Why do we need EDA?

(Behrens & Yu, 2003)

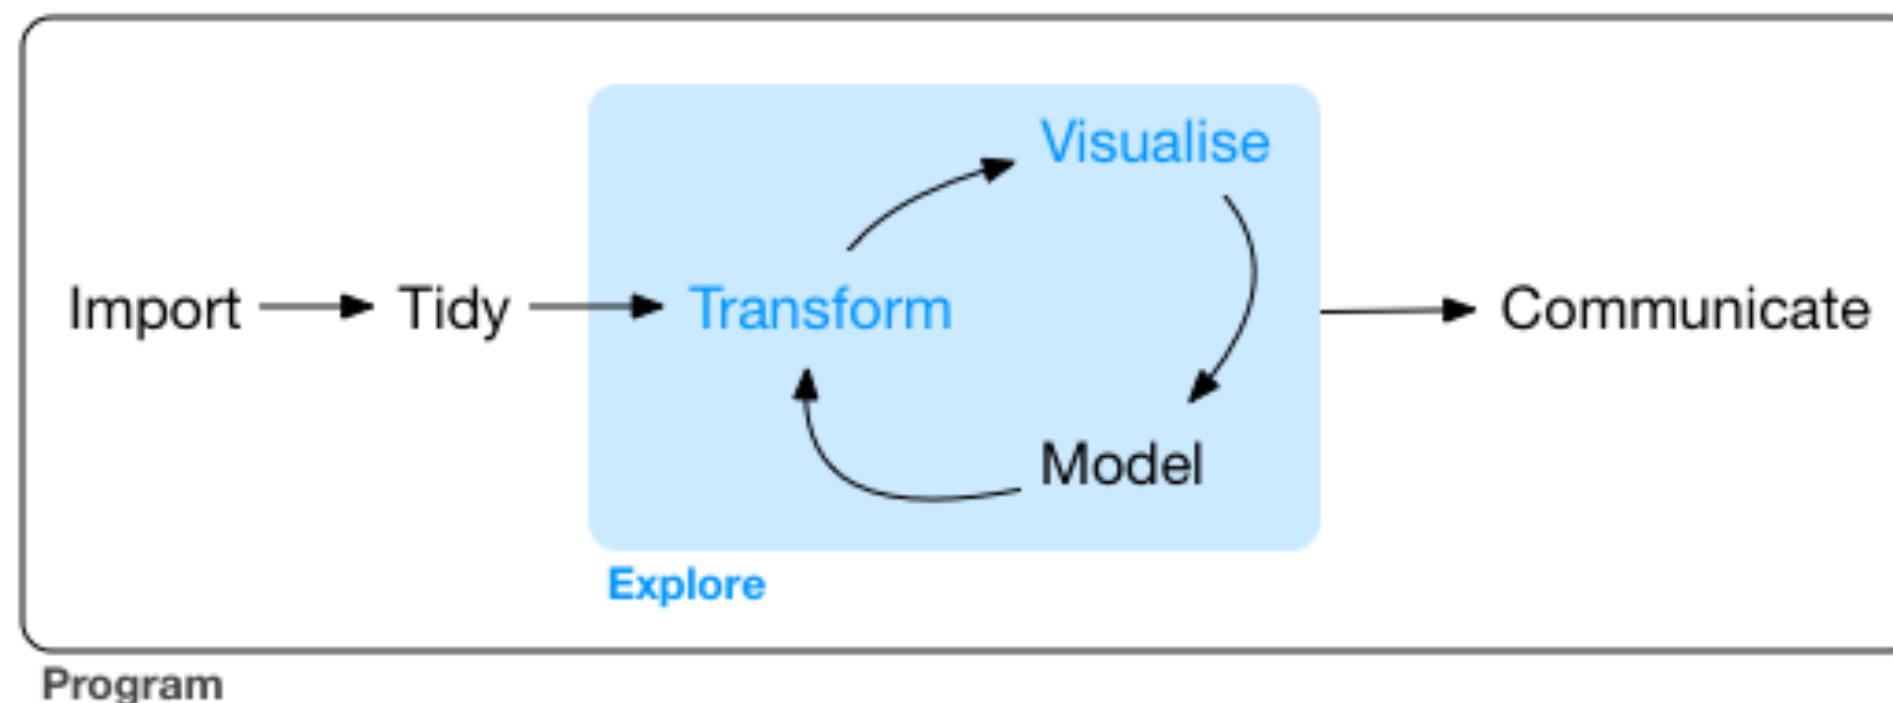
- Summarization = a loss of information
 - If you first look at summarized data (across trials, across participants, etc.), you may miss important patterns that exist at the raw data level
- Statistics lie, so you need graphics
 - Correlations without looking at the scatterplot
 - Means without examining outliers/distribution
 - Statistical tests without examining n

Anscombe's quartet



Where does wrangling stop and EDA begin?

- Data need to be minimally read in, appropriately labelled, and tidied to check and visualize
- EDA will reveal errors or redundancies that will require new data wrangling steps



Tools for data checking that we have already covered

- filter with logical statements

```
> ds %>% filter(class != class_rel)
# A tibble: 34 x 305
   time class class_prop class_rel class_prop_rel x_sum y_sum z_sum corr_xy corr_xz
  <dbl> <fct>    <dbl> <fct>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1  105. held      0.662 sit      0.602  197.  374. -157. -0.0875 0.610
2  414. held      0.657 supine    0.502  216.  31.6   7.50 -0.618 -0.718
3  508. supine    0.896 prone     0.522 -160.  245.  148. -0.771 0.803
4  509. supine    0.647 prone     0.771 -183.  284.   85.7 0.433 -0.0730
5 1065. sit       0.657 prone     0.502 -81.0  249.  440. -0.701 -0.139
```

- fct_count to check factor frequencies

```
> fct_count(ds$class)
# A tibble: 4 x 2
   f         n
  <fct> <int>
1 prone   325
2 held    686
3 sit     812
4 supine  177
```

Quick aside about factors

- When defining factors, all data that do not match a level will be coded as NA
- Checking levels before coding factors (or making a factor without pre-specified levels) should be a first step

```
> ds$class <- factor(ds$class)
> fct_count(ds$class)
# A tibble: 8 x 2
  f           n
  <fct>      <int>
1 held_stat  346
2 held_walk  340
3 prone      325
4 sit_cg      64
5 SIT_CG       1
6 sit_rest   215
7 sit_surf   532
8 supine     177
```

```
> unique(ds$class)
[1] "SIT_CG"      "sit_cg"      "held_stat"  "held_walk"  "supine"     "prone"       "sit_surf"
[8] "sit_rest"
```

Tools for data checking that we have already covered

- summaries (with the right statistics/groupings)

```
> ds_joined %>% summarize(min_age = min(age), max_age = max(age))  
# A tibble: 1 x 2  
  min_age max_age  
  <chr>    <chr>  
1 21      25
```

Tools for data checking that we have already covered

- summaries (with the right statistics/groupings)

```
> ds_joined %>% summarize(min_age = min(age), max_age = max(age))  
# A tibble: 1 x 2  
  min_age max_age  
  <chr>    <chr>  
1 21      25
```

```
# A tibble: 240 x 4  
  participant block condition trial_num  
  <chr>        <chr> <chr>      <dbl>  
1 6191         1    near         1  
2 6191         1    near         2  
3 6191         1    near         3  
4 6191         1    near         4  
5 6191         1    near         5  
6 6191         1    near         6  
7 6191         1    near         7  
8 6191         1    near         8  
9 6191         1    near         9  
10 6191         1    near        10  
# ... with 230 more rows
```

```
> ds %>% group_by(participant, block) %>% summarize(trials_20 = n())  
`summarise()` regrouping output by 'participant' (override with ` .groups ` argument)  
# A tibble: 12 x 3  
# Groups:   participant [2]  
  participant block trials_20  
  <chr>        <chr>    <int>  
1 6191         1         20  
2 6191         2         20  
3 6191         3         20  
4 6191         4         20  
5 6191         5         20  
6 6191         6         20  
7 6192         1         20
```

Tools for data checking that we have already covered

- Automation

- EDA means taking a detailed approach to look at data on different levels (participant/condition/wave/etc.)
- Running multiple filters/checks, plotting multiple figures, etc. can get overwhelming without automation

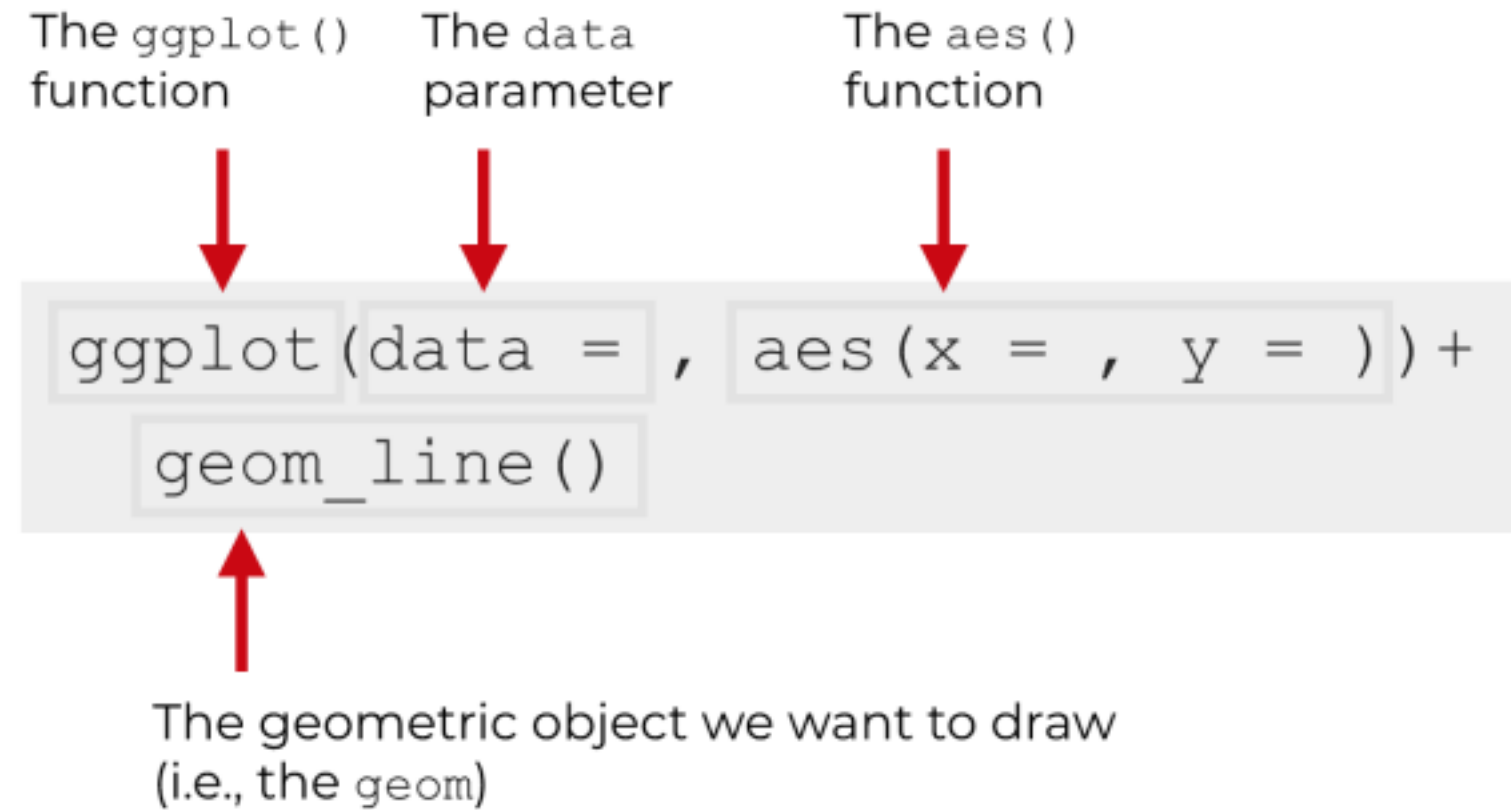
New tools for EDA - Visualizations

- DataExplorer package
 - Brute force, first glance methods
 - plot_histogram() of every continuous variable
 - plot_bar() counts of every categorical variable
- VisDat package
 - vis_miss() to plot missing values
 - vis_expect() to plot conditionals

New tools for EDA - Visualizations


- ggplot2 package (part of tidyverse)
 - Create any type of graph
 - Today we'll talk about making quicker plots for eda using geom_histogram, geom_point, geom_boxplot, and a few others
 - Week 9 we'll talk about making publication-ready plots to communicate effects

Anatomy of a ggplot call




- define the dataset we are using (long format)
- define the mapping of variables to *aesthetics*
- Add (+) geoms, graphical elements like histograms, lines, points, bars, boxplots, and many others
- Optional arguments to change the overall look


Each type of geom has different aesthetics that can be mapped



c + geom_histogram(binwidth = 5) x, y, alpha, color, fill, linetype, size, weight



f + geom_boxplot(), x, y, lower, middle, upper, ymax, ymin, alpha, color, fill, group, linetype, shape, size, weight



e + geom_point(), x, y, alpha, color, fill, shape, size, stroke



h + geom_bin2d(binwidth = c(0.25, 500)) x, y, alpha, color, fill, linetype, size, weight

What aes values are required for each geom?

- Check the help page to see required mappings in bold

Aesthetics

`geom_point()` understands the following aesthetics (required aesthetics are in bold):

- **x**
- **y**
- alpha
- colour
- fill
- group
- shape
- size
- stroke

Adding elements to graphs

- `ggplot() + geomX() +`
- Add (+) other modifications to the plot
 - `xlim(lower_bound, upper_bound)` or `ylim`
 - `hline(yintercept = X)` or `vline`
 - `xlab("x label")`
 - titles, custom scales, other geoms
- Make sure that plus is on the previous line, lines that start with + will throw an error

EDA visualization tutorial

Group work: Exploratory Data Analysis

"259-group-wrangling-checking"

Study setup

- Participants (6 months - 10 years and adults) watched 5 different stimulus videos ("Feist", "Pentatonix", "Science", "Plane", "Dogs") while their eyes were tracked
- The file includes their age in days and their age group (e.g., 0.5-1 year, 1-1.5 years, etc.)
- The DVs are AUC_sal and AUC_dist, which measure how well a saliency vs. a distance model predicted whether participants look (0-1, where 0.5 = chance)
- Precision is a measure of eye tracking data quality, where larger = worse (and over 2.5 is concerning)
- Watched asks whether participants had seen the videos prior to the study

 Filter

	stim	id	age	AUC_sal	AUC_dist	age_group	precision	watched
1	Feist	37	180	0.59957	0.62031	.5-1 y	1.485714	Yes
2	Pentatonix	37	180	0.58893	0.64556	.5-1 y	1.485714	Yes
3	Science	37	180	0.67927	0.82317	.5-1 y	1.485714	No
4	Plane	37	180	0.59959	0.60541	.5-1 y	1.485714	No
5	Dogs	37	180	0.61225	0.54289	.5-1 y	1.485714	No
6	Feist	44	285	0.63152	0.69146	.5-1 y	1.816667	Not Sure
7	Pentatonix	44	285	0.56917	0.75917	.5-1 y	1.816667	Yes
8	Plane	44	285	0.48053	0.63904	.5-1 y	1.816667	No
9	Dogs	44	285	0.61609	0.65665	.5-1 y	1.816667	No
10	Feist	50	240	0.60309	0.64959	.5-1 y	1.616667	No
11	Pentatonix	50	240	0.57844	0.58591	.5-1 y	1.616667	No
12	Science	50	240	0.65497	0.71104	.5-1 y	1.616667	No
13	Plane	50	240	0.60990	0.69281	.5-1 y	1.616667	No
14	Dogs	50	240	0.54050	0.40872	.5-1 y	1.616667	No

Work in groups

- We'll randomly assign you to groups of 3
- Work with your group in RStudio Cloud, and Jake and I will filter in to check on groups
- Do as much as you can, but don't worry if you don't get through all of them
- We will reconvene to go over it together